

# Regression Tasks

Astrini Sie

*asie@seattleu.edu*

**Week 2**

ECEGR4750 - Introduction to Machine Learning  
Seattle University

September 26, 2023

# Recap

- Stable Diffusion demo

## 1 Supervised Learning

- Definition
- Terminologies

## 2 Regression

- Case Study: Input, Model, Cost Function
- Optimization: Ordinary Least Squares (OLS)
- Optimization: Least Mean Squares (LMS)
  - Gradient Descent
  - Batch Gradient Descent
  - Stochastic Gradient Descent

## 3 Case Study, Revisited

# Supervised Learning

## Goal of supervised learning

Given training set with known features and labels, produce a prediction function

# Supervised Learning

## Goal of supervised learning

Given training set with known features and labels, produce a prediction function

- During training: given input data ('**training set**') with **features** and **labels**, learn the relationship ('**prediction function**') between them
- During inference: given a brand new data with **features**, use the **prediction function** to predict the **labels**

# Supervised Learning

## Terminologies:

- Input / **Training Set** consists of feature and label pairs:

Features:  $x^{(1)}, x^{(2)}, \dots, x^{(n)}$ .  $x^{(i)} \in X$ .

Labels:  $y^{(1)}, y^{(2)}, \dots, y^{(n)}$ .  $y^{(i)} \in Y$ .

for  $i = 1, \dots, n$ , where  $n$  is the total number of training samples.

# Supervised Learning

## Terminologies:

- Input / **Training Set** consists of feature and label pairs:

Features:  $x^{(1)}, x^{(2)}, \dots, x^{(n)}$ .  $x^{(i)} \in X$ .

Labels:  $y^{(1)}, y^{(2)}, \dots, y^{(n)}$ .  $y^{(i)} \in Y$ .

for  $i = 1, \dots, n$ , where  $n$  is the total number of training samples.

- Model / **Prediction Function**:

A function  $h_{\theta}(x) : X \rightarrow Y$  that maps the input features  $X$  to the output values  $Y$ , where  $\theta$  is the **parameter** or **weight** of the model. In the training process, we *learn* the values of  $\theta$  that results in good predictions.

# Supervised Learning

## Terminologies:

- Input / **Training Set** consists of feature and label pairs:

Features:  $x^{(1)}, x^{(2)}, \dots, x^{(n)}$ .  $x^{(i)} \in X$ .

Labels:  $y^{(1)}, y^{(2)}, \dots, y^{(n)}$ .  $y^{(i)} \in Y$ .

for  $i = 1, \dots, n$ , where  $n$  is the total number of training samples.

- Model / **Prediction Function**:

A function  $h_{\theta}(x) : X \rightarrow Y$  that maps the input features  $X$  to the output values  $Y$ , where  $\theta$  is the **parameter** or **weight** of the model. In the training process, we *learn* the values of  $\theta$  that results in good predictions.

- Output / **Prediction**:

Given an unseen data point  $x^{(k)}$ , predict the output  $y^{(k)}$  based on prediction function  $h$ .



## Terminologies

- Input / **Training Set** consists of feature and label pairs:  
Features:  $x^{(1)}, x^{(2)}, \dots, x^{(n)}$ .  $x^{(i)} \in X$ .  
Labels:  $y^{(1)}, y^{(2)}, \dots, y^{(n)}$ .  $y^{(i)} \in Y$ .  
for  $i = 1, \dots, n$ , where  $n$  is the total number of training samples.
  - Model / **Prediction Function**:  
A function  $h_{\theta}(x) : X \rightarrow Y$  that maps the input features  $X$  to the output values  $Y$ , where  $\theta$  is the **parameter** or **weight** of the model.  
In the training process, we *learn* the values of  $\theta$  that results in good predictions.
  - Output / **Prediction**:  
Given an unseen data point  $x^{(k)}$ , predict the output  $y^{(k)}$  based on prediction function  $h$ .
- 
- if  $Y$  is continuous, it will be a regression task
  - if  $Y$  is discrete, it will be a classification task

# Supervised Learning

There are two types of supervised learning tasks:

## Regression

if  $Y$  is continuous

- Estimating the relationships between a dependent variable ('**label**') and one or more independent variables ('**features**').
- Example:  $X$  is data of house dimensions and locations, predict  $Y$  the price of the house.

## Classification

if  $Y$  is discrete

- Categorizing a given set of input data into categories ('**classes**') based on one or more variables ('**features**').
- Example:  $X$  is an image, predict if  $Y$  is a "cat" or a "dog".

# Supervised Learning

There are two types of supervised learning tasks:

## Regression

if  $Y$  is continuous

- Estimating the relationships between a dependent variable ('**label**') and one or more independent variables ('**features**').
- Example:  $X$  is data of house dimensions and locations, predict  $Y$  the price of the house.

## Classification

if  $Y$  is discrete

- Categorizing a given set of input data into categories ('**classes**') based on one or more variables ('**features**').
- Example:  $X$  is an image, predict if  $Y$  is a "cat" or a "dog".

# Input

## Regression Sample Case: Predicting House Data

**1) Input:** feature  $x = \text{size}$  and label  $y = \text{price}$ , with  $n = 2$  training samples.

	Size
$x^{(1)}$	2104
$x^{(2)}$	2500
$x^{(3)}$	1600

	Price
$y^{(1)}$	400K
$y^{(2)}$	900K
$y^{(3)}$	330K

# Input

## Regression Sample Case: Predicting House Data

**1) Input:** feature  $x = \text{size}$  and label  $y = \text{price}$ , with  $n = 2$  training samples.

	Size
$x^{(1)}$	2104
$x^{(2)}$	2500
$x^{(3)}$	1600

	Price
$y^{(1)}$	400K
$y^{(2)}$	900K
$y^{(3)}$	330K

**2) Model / Prediction Function:** a linear model

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1$$

# Model

## Regression Sample Case: Predicting House Data

### 2) Model / Prediction Function: a linear model

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1$$

Using the convention of  $x_0 = 1$ , we can rewrite the prediction function as:

$$h(x) = \sum_{i=0}^n \theta_i x_i$$

# Model

## Regression Sample Case: Predicting House Data

### 2) Model / Prediction Function: a linear model

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1$$

Using the convention of  $x_0 = 1$ , we can rewrite the prediction function as:

$$h(x) = \sum_{i=0}^n \theta_i x_i$$

### Goal


Learn the values of  $\theta$  that results in **good prediction**.

**What is a good prediction?**

# Cost Function

## Regression Sample Case: Predicting House Data

### What is a good prediction?



Calculate the distance from each data point (label)  $y^{(i)}$  to the regression line (prediction)  $h_{\theta}(x^{(i)}) = \theta_0 + \theta_1 x_1^{(i)}$ , square it, and sum all of the squared errors together. We call this the **cost function**  $J(\theta)$ .



# Cost Function

## Regression Sample Case: Predicting House Data

**What is a good prediction?** A good prediction is the one that minimizes the cost function  $J(\theta)$ .

### 2a) Cost Function:

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n \left( h_{\theta} \left( x^{(i)} \right) - y^{(i)} \right)^2$$

Choose  $\theta$  such that:

$$\theta = \underset{\theta}{\operatorname{argmin}} J(\theta)$$

**How do we solve for  $\theta$  that minimizes this cost function?**

**How do we solve for this minimization / optimization problem?**

# Solving the Cost Function

Regression Sample Case: Predicting House Data

How do we solve for  $\theta$  that minimizes this cost function?

How do we solve for this minimization / optimization problem?

- ① **Ordinary Least Squares (OLS)**. Analytical solution (closed form).
- ② **Least Mean Squares (LMS)**. Numerical solution (approximation).

# Ordinary Least Squares (OLS)

## Optimization

Choose  $\theta$  that minimizes  $J(\theta)$  using OLS solution. This solution provides a closed form solution of  $\theta$  in terms of the known variables  $x$  and  $y$ .

$$\text{Cost Function: } J(\theta) = \frac{1}{2} \sum_{i=1}^n \left( h_{\theta} \left( x^{(i)} \right) - y^{(i)} \right)^2 \quad (1)$$

We can simplify the expression of the cost function in matrix form using the rule  $z^T z = \sum_i z_i^2$ :

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n \left( h_{\theta} \left( x^{(i)} \right) - y^{(i)} \right)^2 \quad (2)$$

$$J(\theta) = \frac{1}{2} (H - Y)^T (H - Y) \quad (3)$$

# Ordinary Least Squares (OLS)

## Optimization

Let's define  $x$ ,  $y$ , and  $h_{\theta}(x)$  in matrix forms:

$$x^{(1)} = \begin{bmatrix} x_1^{(1)} \\ x_2^{(1)} \\ \vdots \\ x_n^{(1)} \end{bmatrix}, \quad x \in X, \quad X = \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ \vdots \\ (x^{(n)})^T \end{bmatrix}$$

$$y \in Y, \quad Y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(n)} \end{bmatrix}$$

# Ordinary Least Squares (OLS)

## Optimization

Let's define  $x$ ,  $y$ , and  $h_{\theta}(x)$  in matrix forms:

$$h_{\theta}(x) = \sum_{i=0}^n \theta_i x_i$$

In matrix form, this is equivalent to:

$$h_{\theta}(x^{(i)}) = (x^{(i)})^T \theta \quad (4)$$

$$H = X\Theta \quad (5)$$

# Ordinary Least Squares (OLS)

## Optimization

Now, let's plug in  $H$  (5) into the cost function (3):

### Cost Function

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n \left( h_{\theta} \left( x^{(i)} \right) - y^{(i)} \right)^2$$

$$J(\theta) = \frac{1}{2} (H - Y)^T (H - Y)$$

$$J(\theta) = \frac{1}{2} (X\Theta - Y)^T (X\Theta - Y)$$

# Ordinary Least Squares (OLS)

## Optimization

Finally, let's find  $\theta$  that minimizes  $J$ . To do this, find the derivative of  $J$  with respect to  $\theta$ , and set it to zero.

$$\nabla_{\theta} J = \nabla_{\theta} \frac{1}{2} (X\Theta - Y)^T (X\Theta - Y)$$

$$0 = \frac{1}{2} \nabla_{\theta} ((X\Theta)^T - Y^T) (X\Theta - Y)$$

$$0 = \frac{1}{2} \nabla_{\theta} ((X\Theta)^T X\Theta - (X\Theta)^T Y - Y^T (X\Theta) + Y^T Y)$$

Using  $(ab)^T = b^T a^T$  and  $a^T b = b^T a$ :

$$0 = \frac{1}{2} \nabla_{\theta} (\Theta^T X^T X\Theta - Y^T (X\Theta) - Y^T (X\Theta) + Y^T Y)$$

# Ordinary Least Squares (OLS)

## Optimization

continued...

$$0 = \frac{1}{2} \nabla_{\theta} (\Theta^T X^T X \Theta - 2(X^T Y)^T \Theta + Y^T Y)$$

Using  $\nabla_x b^T x = b$  and  $\nabla_x x^T A x = 2Ax$  for symmetric matrix  $A$  :

$$0 = \frac{1}{2} (2X^T X \Theta - 2X^T Y)$$

$$0 = X^T X \Theta - X^T Y$$

Solving for  $\Theta$ :

$$\begin{aligned} X^T X \Theta &= X^T Y \\ \Theta &= (X^T X)^{-1} X^T Y \end{aligned}$$



# Ordinary Least Squares (OLS)

## Optimization

### Ordinary Least Squares (OLS) Solution

The value of  $\theta$  that minimizes  $J(\theta)$  is given in closed form by the equation:

$$\Theta = (X^T X)^{-1} X^T Y$$

# Least Mean Squares (LMS) and Gradient Descent Optimization

Choose  $\theta$  that minimizes  $J(\theta)$  using LMS algorithm. This is a search algorithm in which:

- 1 We start with an “initial guess”  $\theta$
- 2 Repeatedly change  $\theta$  to make  $J(\theta)$  smaller
- 3 Converge to a value of  $\theta$  that minimizes  $J(\theta)$

This is called the **Gradient Descent** algorithm.

# Least Mean Squares (LMS) and Gradient Descent

## Optimization

Choose  $\theta$  that minimizes  $J(\theta)$  using LMS algorithm. This is a search algorithm in which:

① We start with an “initial guess”  $\theta$

② Repeatedly performs an update:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

- $\alpha$  is the **learning rate**
- The update is simultaneously performed  $\forall j = 0, \dots, d$ , where  $j$  is the number of training samples.
- Repeatedly takes a step in the direction of steepest decrease of  $J$ .

③ Converge to a value of  $\theta$  that minimizes  $J(\theta)$

This is called the **Gradient Descent** algorithm.

# Least Mean Squares (LMS) and Gradient Descent Optimization

Least Mean Squares (LMS) / Widrow-Hoff Learning Rule:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Let's solve for this rule by expanding the partial derivative, considering the case of only one training sample  $(x, y)$ , thereby neglecting the sum notation in the cost function  $J(\theta)$ .

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n \left( h_{\theta} \left( x^{(i)} \right) - y^{(i)} \right)^2$$

$$J(\theta) = \frac{1}{2} (h_{\theta}(x) - y)^2$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \frac{1}{2} (h_{\theta}(x) - y)^2$$

# Least Mean Squares (LMS) and Gradient Descent Optimization

continued...

$$\begin{aligned}\frac{\partial}{\partial \theta_j} J(\theta) &= \frac{\partial}{\partial \theta_j} \frac{1}{2} (h_{\theta}(x) - y)^2 \\ &= 2 \cdot \frac{1}{2} (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} (h_{\theta}(x) - y) \\ &= (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} \left( \sum_{i=0}^d \theta_i x_i - y \right) \\ &= (h_{\theta}(x) - y) x_j\end{aligned}$$

Substituting this into the Learning Rule results in:

$$\theta_j := \theta_j - \alpha \left( y^{(i)} - h_{\theta}(x^{(i)}) \right) x_j^{(i)}$$

# Least Mean Squares (LMS) and Gradient Descent Optimization

## Least Mean Squares (LMS) / Widrow-Hoff Learning Rule for one training sample

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$
$$\theta_j := \theta_j - \alpha \left( y^{(i)} - h_{\theta}(x^{(i)}) \right) x_j^{(i)}$$

A couple interesting notes:

- The magnitude of the update is proportional to the error term  $(y^{(i)} - h_{\theta}(x^{(i)}))$ .
- If our prediction  $h_{\theta}(x^{(i)})$  nearly matches the label  $y^{(i)}$ , there is little need to change the parameters  $\theta$ .
- If the prediction  $h_{\theta}(x^{(i)})$  has a larger error from the label  $y^{(i)}$ , then a larger change to the parameters  $\theta$  is needed.

# Batch Gradient Descent

## Optimization

# Stochastic Gradient Descent

## Optimization



# More Complex Data

## Regression Sample Case: Predicting House Data, Revisited

**Input / Training Set:** In the previous example, we only had 1 feature  $x_1 = \text{price}$ . Now we are adding 2 more features  $x_2 = \text{bedrooms}$  and  $x_3 = \text{lotsize}$ .

	Size	Bedrooms	Lot Size
$x^{(1)}$	2104	4	45K
$x^{(2)}$	2500	3	30K

	Price
$y^{(1)}$	400K
$y^{(2)}$	900K

**Model / Prediction Function:** a linear model

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3$$

**Output / Prediction:** Given a new house with size  $x_1^{(k)} = 2250$ , bedrooms  $x_2^{(k)} = 3$ , and lot size  $x_3^{(k)} = 39K$ , predict the price  $y^{(k)}$  of this house.

# References



Chris Re, Andrew Ng, and Tengyu Ma (2023)

CSE229 Machine Learning

*Stanford University*