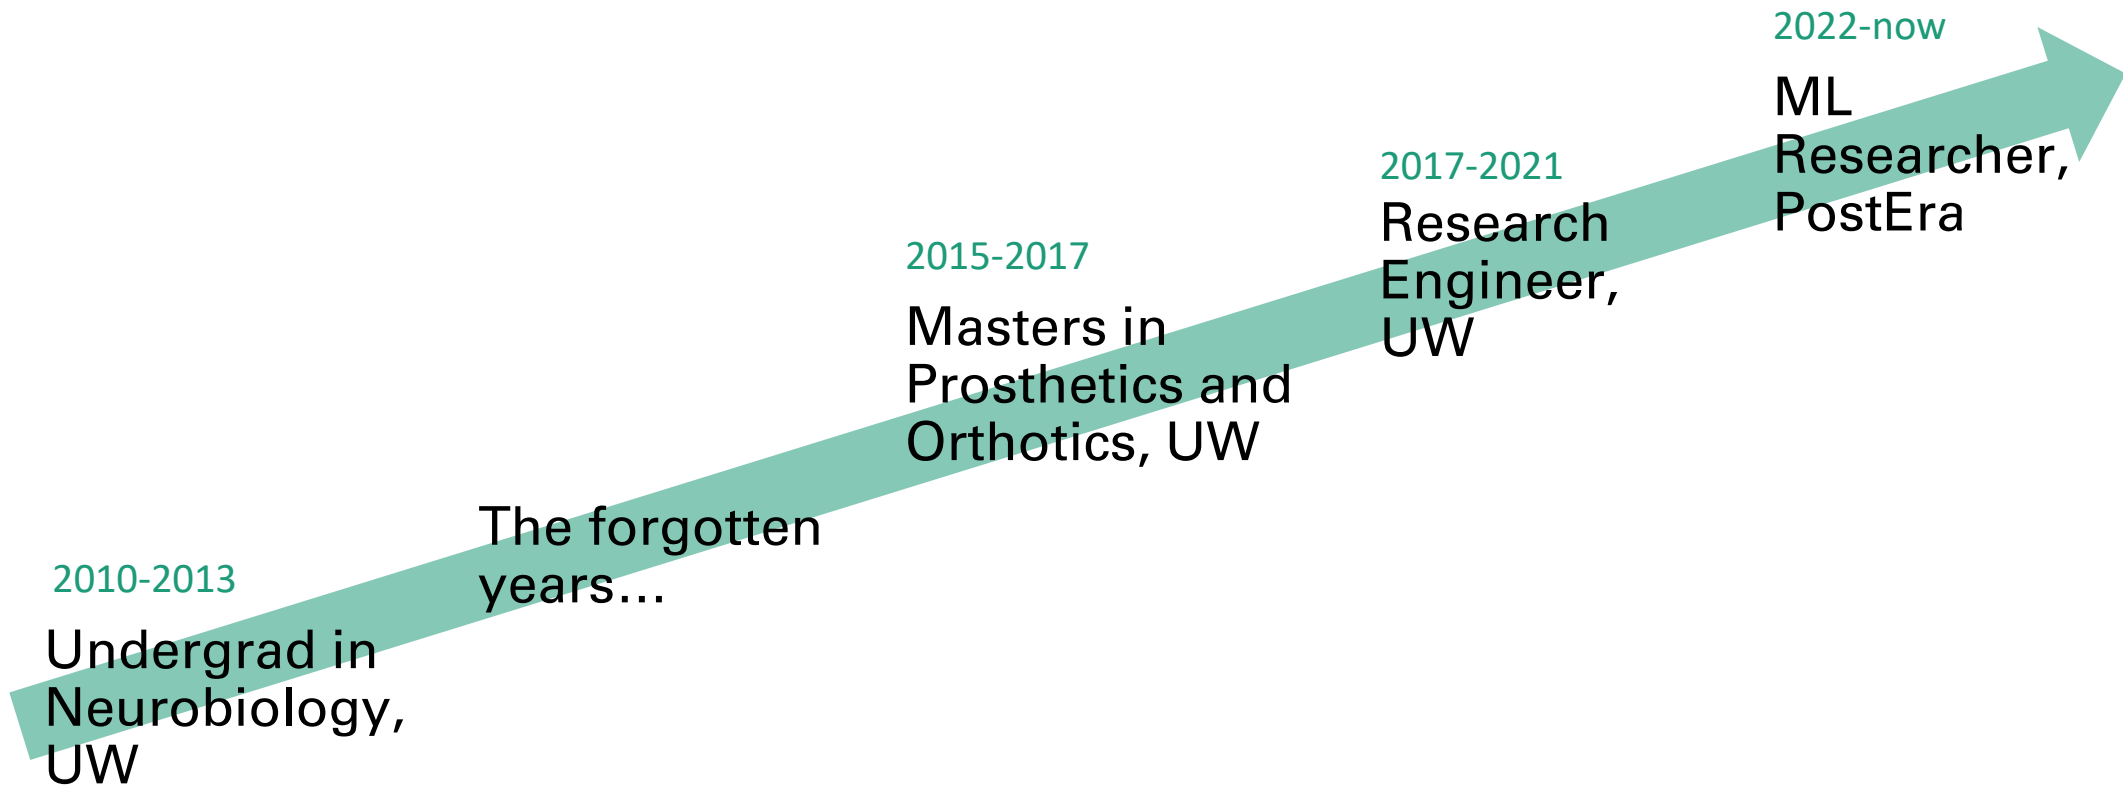




Research Perspectives Machine Learning

david boe - autumn 2023

“typical” career paths for ML



agenda

- History of machine learning research (from a connectionist perspective)
- How to read a research paper
- Good and bad paper presentations to your peers

why does history matter

- Good research is about asking the right questions
- Many breakthroughs in ML are simply applications of old ideas with new technologies
- Gives you a competitive edge
- Informs your ability to search related disciplines for relevant innovations

1943 - perceptron

The “Perceptron” was first thought up

Warren McCulloch,
-philosopher
-psychologist
-neurophysiologist
-45 yrs old



Walter Pitts,
-mathematician
-homeless
-20 yrs old



A LOGICAL CALCULUS OF THE IDEAS IMMANENT IN NERVOUS ACTIVITY*

■ WARREN S. MCCULLOCH AND WALTER PITTS
University of Illinois, College of Medicine,
Department of Psychiatry at the Illinois Neuropsychiatric Institute,
University of Chicago, Chicago, U.S.A.

Because of the “all-or-none” character of nervous activity, neural events and the relations among them can be treated by means of propositional logic. It is found that the behavior of every net can be described in these terms, with the addition of more complicated logical means for nets containing circles; and that for any logical expression satisfying certain conditions, one can find a net behaving in the fashion it describes. It is shown that many particular choices among possible neurophysiological assumptions are equivalent, in the sense that for every net behaving under one assumption, there exists another net which behaves under the other and gives the same results, although perhaps not in the same time. Various applications of the calculus are discussed.

1. Introduction. Theoretical neurophysiology rests on certain cardinal assumptions. The nervous system is a net of neurons, each having a soma and an axon. Their adjunctions, or synapses, are always between the axon of one neuron and the soma of another. At any instant a neuron has some threshold, which excitation must exceed to initiate an impulse. This, except for the fact and the time of its occurrence, is determined by the neuron, not by the excitation. From the point of excitation the impulse is propagated to all parts of the neuron. The velocity along the axon varies directly with its diameter, from $< 1 \text{ ms}^{-1}$ in thin axons, which are usually short, to $> 150 \text{ ms}^{-1}$ in thick axons, which are usually long. The time for axonal conduction is consequently of little importance in determining the time of arrival of impulses at points unequally remote from the same source. Excitation across synapses occurs predominantly from axonal terminations to somata. It is still a moot point whether this depends upon irreducibility of individual synapses or merely upon prevalent anatomical configurations. To suppose the latter requires no hypothesis *ad hoc* and explains known exceptions, but any assumption as to cause is compatible with the calculus to come. No case is known in which excitation through a single synapse has elicited a nervous impulse in any neuron, whereas any neuron may be excited by impulses arriving at a sufficient number of neighboring synapses within the period of latent addition, which lasts $< 0.25 \text{ ms}$. Observed temporal summation of impulses at greater intervals

* Reprinted from the *Bulletin of Mathematical Biophysics*, Vol. 5, pp. 115-133 (1943).


99

LITERATURE

Carnap, R. 1938. *The Logical Syntax of Language*. New York: Harcourt-Brace.
Hilbert, D. and W. Ackermann. 1927. *Grundzüge der Theoretischen Logik*. Berlin: Springer.
Russell, B. and A. N. Whitehead. 1925. *Principia Mathematica*. Cambridge University Press.

1943 - perceptron

Inspired by physiology of the brain



Cajal, 1904,
neuroanatomist
and artist

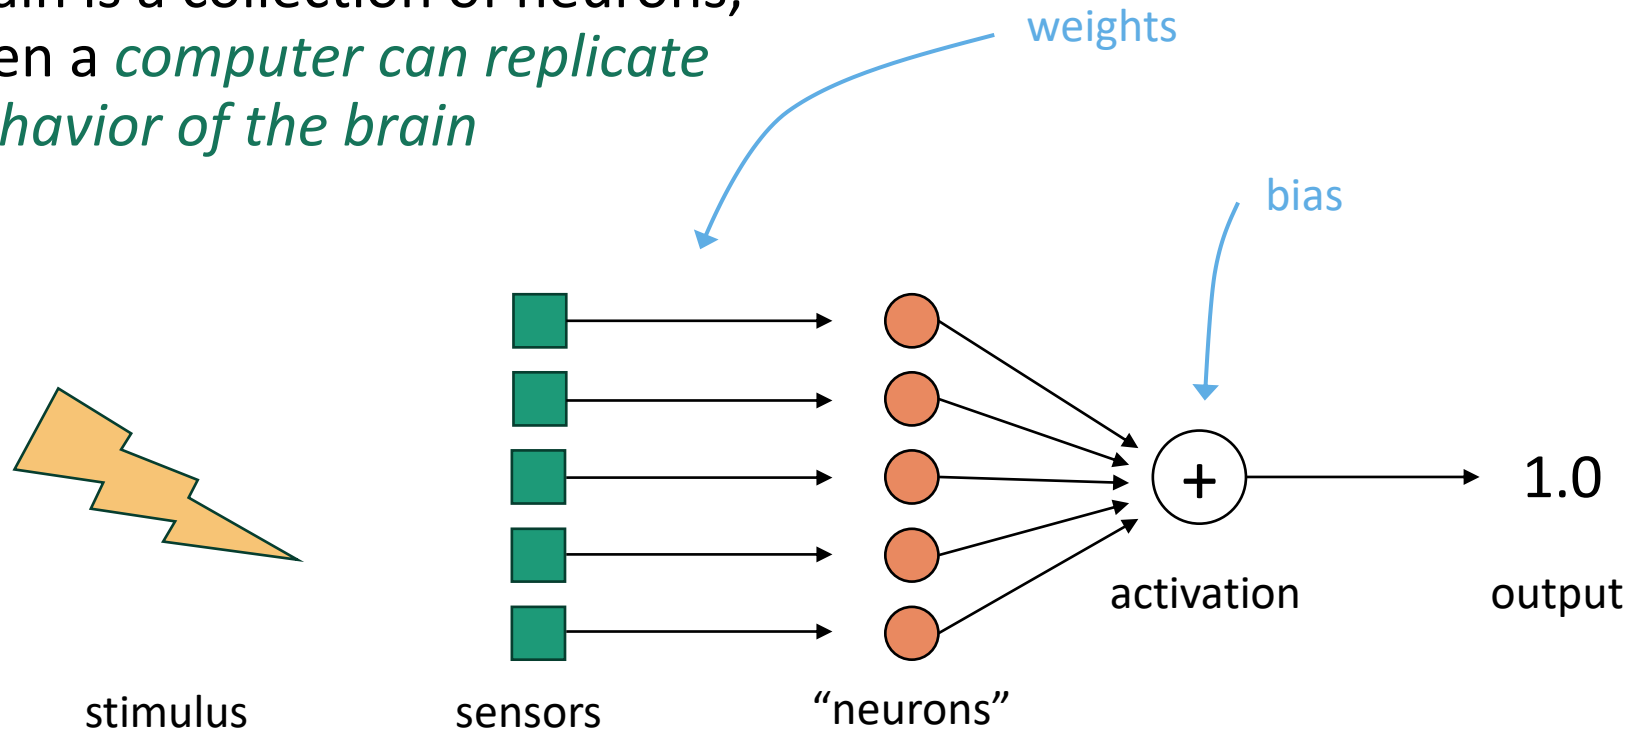
Every dendrite has a “weight”

Dendrites excite the neuron proportional to their weight

If the neuron excitation rises above a threshold, then the neuron “fires”

1943 - perceptron

If a computer could replicate behavior of a neuron, and the brain is a collection of neurons, then a *computer can replicate behavior of the brain*



1943 - perceptron

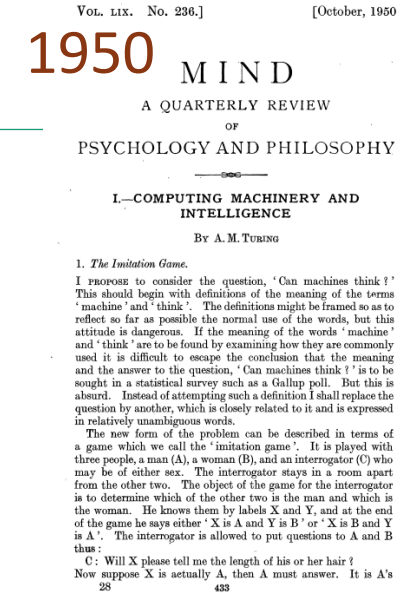
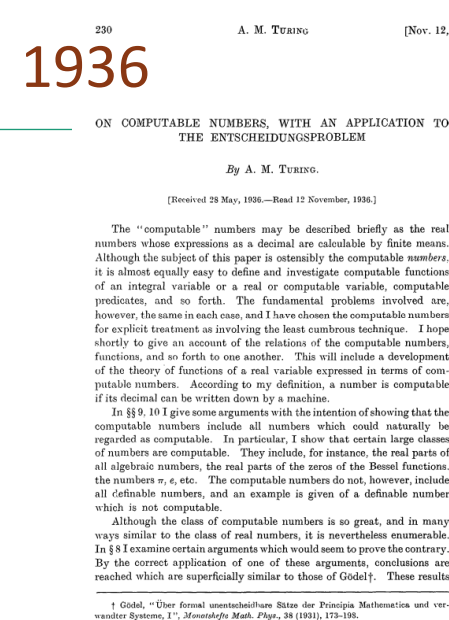
Formalized the concept of *training* (not *programming*) a computer. Idea that computers could *learn* on their own was revolutionary.

Gave rise to the field of neural networks, or connectionism


meanwhile...

Alan Turing is laying the theoretical foundation for computing, launching the field of artificial intelligence.

“Symbolic AI” rises as the dominant research paradigm throughout the 1940s and 1950s.



1958 – perceptron v1



Frank Rosenblatt

-cognitive scientist

-hype man


“The Navy revealed the embryo of an electronic computer today that it expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence . . . Later Perceptrons will be able to recognize people and call out their names and instantly translate speech in one language to speech and writing in another language, it was predicted. (New York Times, 1958a, p. 25 :2)”

1958 – perceptron v1

Rosenblatt's exuberant personality + sensational media reporting made him the enemy of symbolic AI.

Enter his nemesis, Marvin Minsky, mathematician.

Minsky built his own hardware implementation of a neural network in 1951, called SNARC (Stochastic Neural Analog Reinforcement Calculator).



1958 – 1965

Minsky and Rosenblatt were both well aware of the challenges of neural networks. To name a few...


- There was no way to know if a perceptron could solve a problem *a priori*
- Could add layers, but nobody figured out how to adjust the weights
- Networks had poor generalization
- Not SE3 invariant
- Could not solve simple problems like XOR
- Could only work on linearly separable classes

1960s – the Minsky takedown

Marvin Minsky and Seymour Papert (also mathematician) published *Perceptrons* in 1969.

They set out to describe the shortcomings of the connectionist movement, once and for all (and divert funding towards their own work).

Connectionism was increasingly viewed with skepticism, and this book fully killed it off.



1970s – first AI winter

Connectionism failure to deliver brought down the whole AI field. Analog computers were discarded. Funding dried up. Scientists moved on to other areas.

But machine learning scientists quietly continued...

During this time, several new methods were born

- Digital computing gained power
- Least mean squares (LMS) – *Widrow and Hoff*
- Neocognitron for recognizing Japanese handwriting, *Kunihiko Fukushima*
- First self-organizing networks, *Shun'ichi Amari, William Little*
- Automatic differentiation, *Seppo Linnainmaa*

what did we learn the first time around?

Overhyping expectations can tank an entire scientific discipline (don't cross mathematicians)

Translation from the lab (AI that plays checkers) to the real world (AI that does taxes) is non-trivial

Machine learning can draw assistance from many other fields

- Robotics
- Physics
- Computing
- Linguistics
- Mathematics
- Cognition
- ...?

machine learning needs a savior


Symbolic AI work was hitting its ceiling. Clunky computers and man-made schemas couldn't solve longstanding problems like speech recognition or machine translation.

Then, a physicist to the rescue!

John Hopfield (physicist, chemist, biologist)
comes up with the next big thing.



1982 – hopfield networks



1982 – hopfield networks


Like the Perceptron, Hopfield networks learn to recognize patterns. Unlike the Perceptron, they can still work in noisy conditions.

Operates over a fully connected graph, with symmetric edges.

Each node is connected to one input.

Learns the weights (edges) with Hebb's rule.


Prediction happens recursively.



1982 – hopfield networks

For images, each pixel is assigned a node.

Weights between pixels i, j that have the same state are increased in strength (and decreased when pixels are different). Do this for every pattern p .



$$w_{ij} = \sum_p x_i^{(p)} x_j^{(p)}$$


$$w_{ij} = 1.0 * 1.0$$

$$w_{ij} = 1.0 * -1.0$$


1982 – hopfield networks

If we teach a Hopfield network to learn zeros, here is what the weight matrix looks like

For neuron 1





For neuron 4



1982 – hopfield networks

If we feed a Hopfield network an 8, it will still converge to a stable energy state, a zero.

Viewed as a dynamical system, the zero state has become an “attractor”



Remind you of
anything?

1982 – hopfield networks

Re-ignited interest in “connectionist” ideas as physicists, engineers, neuroscientists, etc. all saw their own ideas reflected in the algorithm.

Hebb’s rule – classic rule by which neurons learn

Energy based functions – typical way that physicists think of stateful problems

Content-addressable-memory – a sought after feature for building memory systems

Was one of the first “recurrent” neural networks, and was “unsupervised”

1982 – hopfield networks

Unfortunately, there was no killer application. Wasn't a regressor, wasn't a classifier. Was something different.

Capacity issues (couldn't learn very many patterns) combined with 1980s computing power meant there was a real ceiling on what they could do.

“Fully connected graph” architecture was not well suited to problems of the day like image recognition and time series forecasting.

1986 – backprop to the future

Scientists revisited the old set of perceptron challenges, namely, how do you update the weights of multi-layer perceptrons?

Inertia built across multiple people and institutions.

David Rumelhart (psychologist),
Geoffrey Hinton (psych, comp neuro, CS),
Ronald Williams (cog sci?)
published this landmark paper.

Learning representations by back-propagating errors

David E. Rumelhart*, Geoffrey E. Hinton†
& Ronald J. Williams*

* Institute for Cognitive Science, C-015, University of California,
San Diego, La Jolla, California 92093, USA

† Department of Computer Science, Carnegie-Mellon University,
Pittsburgh, Philadelphia 15213, USA

We describe a new learning procedure, back-propagation, for networks of neurone-like units. The procedure repeatedly adjusts the weights of the connections in the network so as to minimize a measure of the difference between the actual output vector of the net and the desired output vector. As a result of the weight adjustments, internal 'hidden' units which are not part of the input or output come to represent important features of the task domain, and the regularities in the task are captured by the interactions of these units. The ability to create useful new features distinguishes back-propagation from earlier, simpler methods such as the perceptron-convergence procedure¹.

There have been many attempts to design self-organizing neural networks. The aim is to find a powerful synaptic modification rule that will allow an arbitrarily connected neural network to develop an internal structure that is appropriate for a particular task domain. The task is specified by giving the desired state vector of the output units for each state vector of the input units. If the input units are directly connected to the output units it is relatively easy to find learning rules that iteratively adjust the relative strengths of the connections so as to progressively reduce the difference between the actual and desired output vectors². Learning becomes more interesting but

more difficult when we introduce hidden units whose actual or desired states are not specified by the task. (In perceptrons, there are 'feature analysers' between the input and output that are not true hidden units because their input connections are fixed by hand, so their states are completely determined by the input vector: they do not learn representations.) The learning procedure must decide under what circumstances the hidden units should be active in order to help achieve the desired input-output behaviour. This amounts to deciding what these units should represent. We demonstrate that a general purpose and relatively simple procedure is powerful enough to construct appropriate internal representations.

The simplest form of the learning procedure is for layered networks which have a layer of input units at the bottom; any number of intermediate layers; and a layer of output units at the top. Connections within a layer or from higher to lower layers are forbidden, but connections can skip intermediate layers. An input vector is presented to the network by setting the states of the input units. Then the states of the units in each layer are determined by applying equations (1) and (2) to the connections coming from lower layers. All units within a layer have their states set in parallel, but different layers have their states set sequentially, starting at the bottom and working upwards until the states of the output units are determined.

The total input, x_j , to unit j is a linear function of the outputs, y_i , of the units that are connected to j and of the weights, w_{ji} , on these connections

$$x_j = \sum_i y_i w_{ji} \quad (1)$$

Units can be given biases by introducing an extra input to each unit which always has a value of 1. The weight on this extra input is called the bias and is equivalent to a threshold of the opposite sign. It can be treated just like the other weights.

A unit has a real-valued output, y_j , which is a non-linear function of its total input

$$y_j = \frac{1}{1 + e^{-x_j}} \quad (2)$$

† To whom correspondence should be addressed.

1986 – backprop to the future

Suddenly, multi-layer perceptrons were all the rage again.

Backpropagation was widely applicable to many kinds of models, encouraging a new wave of research.

Apply the chain rule to partial derivatives to find each weight's gradient with respect to the error.

1988 – convolutional neural networks

Hubel and Wiesel's famous experiments in 1959 on the visual cortex inspired the Neocognitron in 1979.

Now equipped with backpropagation, the CNN was born.

Proceedings of annual conference of the Japan Society of Applied Physics, Sept. 1988

6 p-M-13

光アソシアトロン (V) —学習の構造とその改良—
OPTICAL ASSOCIATRON —Structure of Learning and Its Improvement—
見沼科学研究所 石川正俊 浜松ホトニクス株式会社 豊田晴義 向坂直久 鈴木義二
[PR] N. Ishikawa HPK H. Toyoda N. Mukozakura Y. Suzuki

[illegible]


Wei Zhang, 1988

Alphabet detection

Medical imaging

6 p-M-14

シフト・インバリエント型パターン認識ニューラル・ネットワークとその光アーキテクチャー
Shift invariant pattern recognition neural network and its optical architecture
阪大工 雄偉, 谷田 純, 伊東 一良, 岡 芳樹

[illegible]

④1. Neural Network Model.
Chart, et al., Parallel Distrib
2,

Communicated by Dana Ballard

Backpropagation Applied to Handwritten Zip Code Recognition

Y. LeCun
B. Boser
J. S. Denker
D. Henderson
R. E. Howard
W. Hubbard
L. D. Jackel
AT&T Bell Labs

Yann LeCun, 1989

The ability of learning networks to generalize can be greatly enhanced by providing constraints from the task domain. This paper demonstrates how such constraints can be integrated into a backpropagation network through the architecture of the network. This approach has been successfully applied to the recognition of handwritten zip code digits provided by the U.S. Postal Service. A single network learns the entire recognition operation, going from the normalized image of the character to the final classification.

1 Introduction

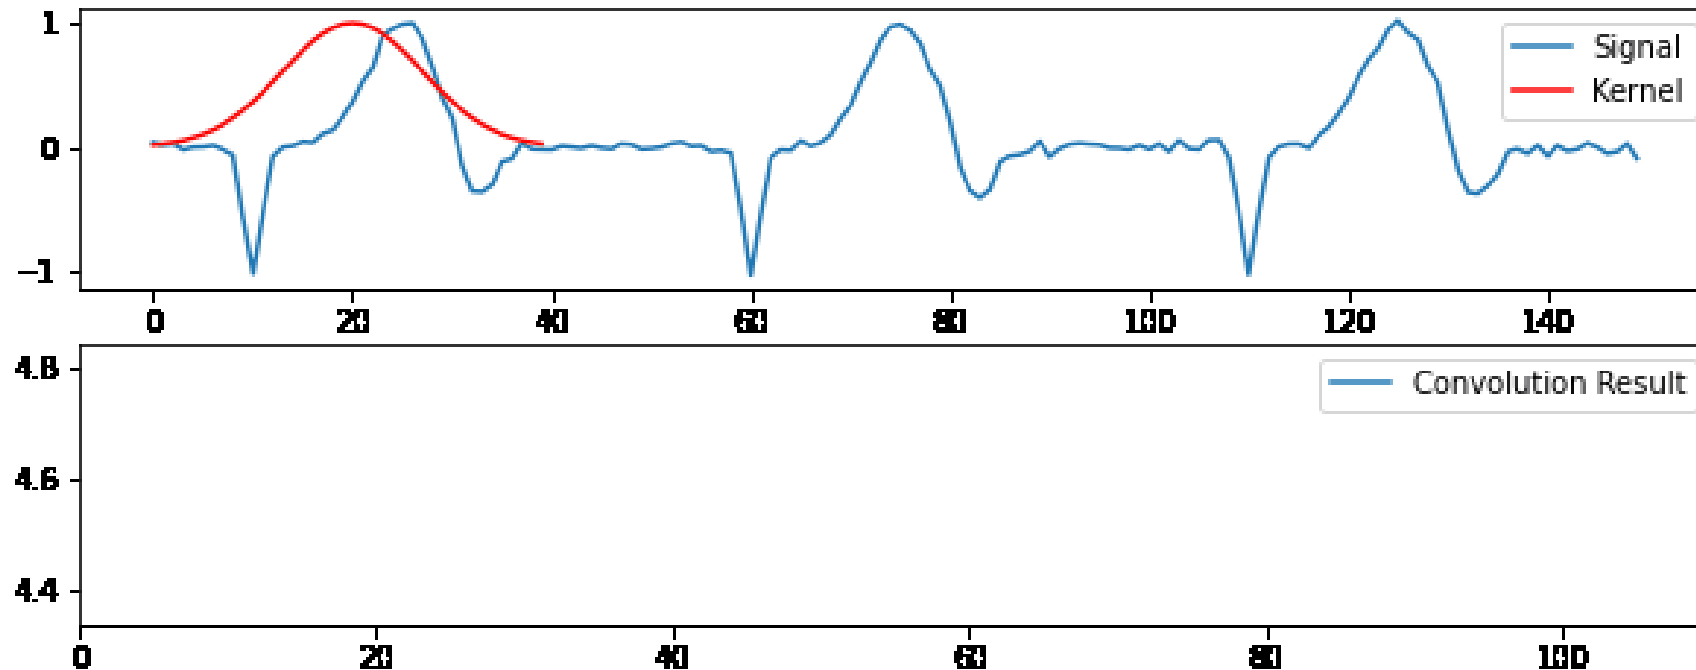
Previous work performed on recognizing simple digit images (LeCun 1989) showed that good generalization on complex tasks can be obtained by designing a network architecture that contains a certain amount of a priori knowledge about the task. The basic design principle is to reduce the number of free parameters in the network as much as possible without overly reducing its computational power. Application of this principle increases the probability of correct generalization because it results in a specialized network architecture that has a reduced capacity (Denski et al. 1987; Pataczny et al. 1987; Tiedey et al. 1989; Cui et al. 1989), and a reduced Vapnik-Chervonenkis dimensionality (Baum and Hasselauer 1989).

In this paper, we apply the backpropagation algorithm (Rumelhart *et al.* 1986) to a real-world problem in recognizing handwritten digits taken from the US Mail. Unlike previous results reported by our group or this problem (Denker *et al.* 1989), the learning network is directly fed with images, rather than feature vectors, thus demonstrating the ability of backpropagation networks to deal with large amounts of low-level information.

Optical character recognition (OCR)

1988 – convolutional neural networks



A convolution is when a signal is convolved with a *kernel* or *filter*




1988 – convolutional neural networks

A convolutional neural net simply *learns* those kernels, to solve some task.

- Translation invariant
- Parameter sharing + pooling = efficiency on images

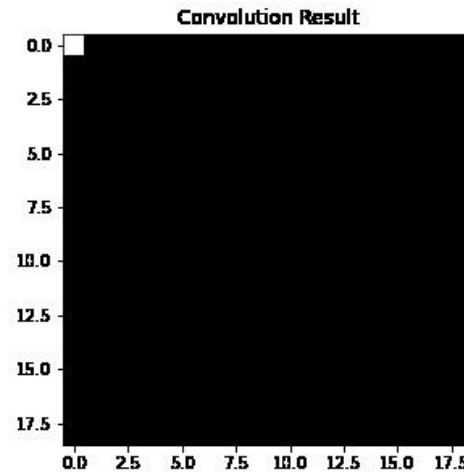




Edge kernel






RBF kernel

1988 – convolutional neural networks



Can also stack convolutions on each other.

CNNs tend to learn primitive features like edges first, then learn more abstract features like “cat” in deeper layers.



Hubel and Wiesel showed that the visual cortex operates in much the same way, 30+ years prior!

1990s – here we go again

Demand for AI hardware drove innovations in massively parallel computers.

US and Japan competed to make the most powerful AI computing systems.

Symbolic AI was the dominant paradigm, and was making its move on corporate America.



1990s – the second AI winter

Much like the first AI winter, expectations exceeded reality. No real world market for AI.

Symbolic AI / expert systems couldn't handle real world data.

Connectionism was incredibly expensive and lacked power.

Bayesian statisticians got their time to shine.



1990s – the second AI winter

Symbolic AI was largely abandoned by researchers.

Connectionism research continued on, albeit quietly. Several notable advances were

- Q-learning in reinforcement learning
- Support Vector Machines
- Long Short Term Memory networks

Probability and statistics exerted their influence on ML

2000s – the rise of big data

By 1996, storing data digitally was cheaper than on paper.

The internet grew up, and data became a commodity for online companies.

Still, machine learning research remained in the academic sphere.

2012 – AlexNet

Alex Krizhevsky and Ilya Sutskever entered the ImageNet competition.


SVMs, segmentation methods, kernel methods, etc. were winners of prior years competitions.

They wanted to use a many-layered neural network, but training it on a million images would have taken months. So they used a pair of Nvidia GTX 580s, the current best parallel processing hardware available.

2012 – AlexNet

AlexNet won the competition handedly, with a top-5 error of 0.15. The next nearest was 0.26!

Amazingly, AlexNet had 62 *million* parameters over 8 layers. Google snatched them right up, and the term “deep learning” was coined.



2012 – 2016 – industry adoption

Tech companies all jumped on to the “deep learning” bandwagon.

Access to GPU clusters and large datasets put industrial research on equal footing with academics.


Finally, machine learning found its market fit.

2014 – generative AI revival

Generative Adversarial Networks, first described in 1991, were revived by Ian Goodfellow.



Variational Autoencoders, an application of Bayesian theory to deep learning, also revived.



2017 – _____ is all you need

Transformer paper published by Google Brain.

Transformers dominated for autoregressive and sequence analysis tasks, like translation.

Transformer-style models have roots in the 90s. The “attention” method was built on prior work in RNNs. Merging the two methods on large datasets resulted in impressive results.


One problem though, the transformer architecture is data hungry and computationally demanding.

also started a super obnoxious naming trend

2018 to today – the diaspora

Machine learning became a more accessible field of research to enter

- arXiv made research available for everyone across the world
- availability of supplementary code with publications
- development of ML libraries like pytorch and tensorflow
- open source approach overran gatekeeping by institutions



2018 to today – GPT

GPT (generative pre-trained transformer) models are essentially transformers, with the bells and whistles that come with big data pretraining tasks.

Model	Architecture	Parameter count	Training data
GPT-1	12-level, 12-headed Transformer decoder (no encoder), followed by linear-softmax.	117 million	BookCorpus . ^[27] 4.5 GB of text, from 7000 unpublished books of various genres.
GPT-2	GPT-1, but with modified normalization	1.5 billion	WebText: 40 GB of text, 8 million documents, from 45 million webpages upvoted on Reddit .
GPT-3	GPT-2, but with modification to allow larger scaling	175 billion ^[31]	499 Billion tokens consisting of CommonCrawl (570 GB), WebText, English Wikipedia, and two books corpora (Books1 and Books2).
GPT-3.5	Undisclosed	175 billion ^[31]	Undisclosed
GPT-4	Also trained with both text prediction and RLHF ; accepts both text and images as input. Further details are not public. ^[26]	Undisclosed	Undisclosed

2022 to today – Stable Diffusion

Diffusion models have their roots in earlier frameworks, like Hopfield Networks.

Stable Diffusion was first to train such a model on billions of images. Beats GANs in image generation capabilities.

Generative techniques are widely applicable to other domains like audio, video, 3d renderings

research landscape today

Research is highly fragmented into specialties by industry and model architectures.

Computational chemistry uses graph neural networks.

Self-driving cars use computer vision.

LLM research uses large pretrained models with finetuning methods.

Financial industry uses uncertainty modeling and time series methods

And so on and so forth.

research landscape today

The accessibility of machine learning research comes with risks

- Focus on public benchmarks incentivizes chasing metrics for glory
- Although code is often provided, it is rarely plug-and-play
- Most papers have novelty that is domain specific or narrow in scope
- Quality of research as a whole suffers without peer review

reading machine learning papers

Some tips for reading research papers.

I do most of my literature searches with simple google searches. Finding the right keywords can sometimes be nontrivial.

Let's say I want to understand more about Graph Neural Networks. This paper comes up: <https://arxiv.org/abs/2306.03589>

step 1: authors, institutions, publication

Before you do anything, check the authors and the institutions. Have you heard of them? Do you trust them?

Francesco Di Giovanni*
University of Cambridge
fd405@cam.ac.uk

T. Konstantin Rusch*
ETH Zürich
konstantin.rusch@sam.math.ethz.ch

Michael M. Bronstein
University of Oxford

Andreea Deac
Mila, Université de Montréal

Marc Lackenby
University of Oxford

Siddhartha Mishra
ETH Zürich

Petar Veličković
Google DeepMind

Yes! These are all heavy hitters in ML research.

But, this paper is in arXiv, which is not peer-reviewed. This means you should be very critical, regardless of the authors.

step 2: read the abstract


Read the abstract. You will probably not understand everything at first, but if the paper is well written, it will tell you if it is relevant to your interests.

Abstract

Graph Neural Networks (GNNs) are the state-of-the-art model for machine learning on graph-structured data. The most popular class of GNNs operate by exchanging information between adjacent nodes, and are known as Message Passing Neural Networks (MPNNs). Given their widespread use, understanding the expressive power of MPNNs is a key question. However, existing results typically consider settings with uninformative node features. In this paper, we provide a rigorous analysis to determine which function classes of node features can be learned by an MPNN of a given capacity. We do so by measuring the level of *pairwise interactions* between nodes that MPNNs allow for. This measure provides a novel quantitative characterization of the so-called over-squashing effect, which is observed to occur when a large volume of messages is aggregated into fixed-size vectors. Using our measure, we prove that, to guarantee sufficient communication between pairs of nodes, the capacity of the MPNN must be large enough, depending on properties of the input graph structure, such as commute times. For many relevant scenarios, our analysis results in impossibility statements in practice, showing that over-squashing hinders the expressive power of MPNNs. We validate our theoretical findings through extensive controlled experiments and ablation studies.


step 3: skim the figures

Do NOT read the introduction next! Instead, skim through the figures and captions of the paper. A good paper worth your read conveys its main message just in the figures.




Modeling the relationship between nodes with springs is novel and interesting!

$\tau(\bullet, \bullet) \approx 30$
 $\tau(\bullet, \bullet) \approx 36$
 $\tau(\bullet, \bullet) \approx 252$



Commute time is a new term to me, but it seems important here, so I should pay attention to that.



I'm gathering the main takeaway of this paper is really that GNNs need to have sufficient capacity in order to realize their full potential. Very relevant!

step 4: lightly read the introduction

Depending on your knowledge of that research domain, you could skim through to the *end* of the introduction. That is where authors describe the *novelty* of their method.

If you are reading to learn more about the domain, then read the introduction line by line. Go look up the references for everything you don't understand. Skim those papers to get an understanding of why they matter.

step 4: lightly read the introduction

Contributions. Our main goal is to address the above question and show how *over-squashing* can be understood as the *misalignment between the task and the graph-topology*, ultimately *limiting the classes of functions that MPNNs of practical size can learn* (see **Figure 1**). We start by measuring the extent to which an MPNN allows any pair of nodes in the graph to *interact* (via *mixing* their features). With this measure as a tool, we characterize which functions of node features can be learned by an MPNN and how the model architecture and parameters, as well as the topology of the graph, affect the expressive power. More concretely,

Authors did a good job at calling out their specific contributions. Always look for this!

Every citation is a full paper that cannot be summarized in a single sentence. It is common practice to set up a strawman with citations to make your own work look more impactful. Ask yourself – did this author accurately interpret those citations?

step 5: lightly read the rest of the paper

You've read the introduction, you know what the authors are trying to do. Read the rest of the paper, but don't get hung up on things you don't understand quite yet.

$$\mathbf{h}_v^{(t)} = f^{(t)}\left(\mathbf{h}_v^{(t-1)}, g^{(t)}\left(\left\{\!\!\left\{\mathbf{h}_u^{(t-1)}, (v, u) \in \mathbb{E}\right\}\!\!\right\}\right)\right), \quad \mathbf{h}_v^{(0)} = \mathbf{x}_v \quad (1)$$

where $f^{(t)}, g^{(t)}$ are learnable functions and the aggregation function $g^{(t)}$ is invariant to permutations. Specifically, we study a class of MPNNs of the following form,

$$\mathbf{h}_v^{(t)} = \sigma\left(\mathbf{\Omega}^{(t)}\mathbf{h}_v^{(t-1)} + \mathbf{W}^{(t)} \sum_u \mathbf{A}_{vu} \psi^{(t)}(\mathbf{h}_v^{(t-1)}, \mathbf{h}_u^{(t-1)})\right), \quad \mathbf{h}_v^{(0)} = \mathbf{x}_v, \quad (2)$$

80% of the time this level of notation is unnecessary for understanding an ML paper. Don't bother reading it if it is rehashing old concepts or proofs. You can always come back to it later.

step 5: lightly read the rest of the paper

Often, the methods will be pretty obtuse. Don't worry, it's not you, it's them. Rarely is a paper written so it can be understood easily.

For this paper, I read the section headers, and decided to skip to Section 4.

4 Over-squashing limits the expressive power of MPNNs

After skimming it, it seems they make their own metric. When the metric is large, the model is over-squashing. When it's small, the model is expressing. Carry on!

They are interested in testing it in two scenarios – making models deeper or making models wider.

step 5: lightly read the rest of the paper


5 Experimental validation of the theoretical results

Experimental results time! If these don't exist, I may not even read the paper. Deducing the value of the paper's findings is difficult without some kind of experiment.

Here I am usually looking for a few things

- What is the dataset used
- What are the model architectures, training tasks, and loss functions
- What are the comparisons being made

I will make a note on the key notations used, so I can go back and figure them out.



step 6: read the conclusion

Nowadays, ML papers are not typically divided into traditional categories. Still, I want to read the authors' summary of their paper. If I look at the end, I will often find this, along with a description of the limitations of their approach.

If there is no limitations section, feel free to boo.

In this paper, it is all inside the Discussion section. The discussion usually rehashes the introduction, and then adds the authors interpretation of the results in the context of their cited works.

step 7: (optional) read it all again, slower

You should already have a good idea of the central claims, the methods, and a sense of how much you believe in it. Often, that is enough to help you on your way (feel free to cite it now!).

If your motivation is deeper, now is when you go back and learn all that notation you missed. Don't feel compelled to understand it *all*.

Write down or discuss your thoughts on the paper with your colleagues, everyone catches something different. You can learn it *much* faster this way.

your research presentation

You will present a research paper to your classmates.

You will have 15 minutes to present, followed by a 5 min Q&A.

You will give and receive constructive evaluations for your peers.

Deep Residual Learning for Image Recognition

[2016 IEEE Conference on Computer Vision and Pattern Recognition \(CVPR\)](#)

Kaiming He

Xiangyu Zhang

Shaoqing Ren

Jian Sun

Microsoft Research

{kahe, v-xiangz, v-shren, jiansun}@microsoft.com

Introduction

- This paper is about deep residual learning for image recognition
- This technique got 1st place in ImageNet and COCO tasks
- Residual connections are also called shortcuts




Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

Related Work

VLAD, Fisher Vectors, and Multigrid methods have also been used on residual connections

Lots of papers have studied shortcuts and tried many things like gated shortcuts

In this paper, shortcuts are always open

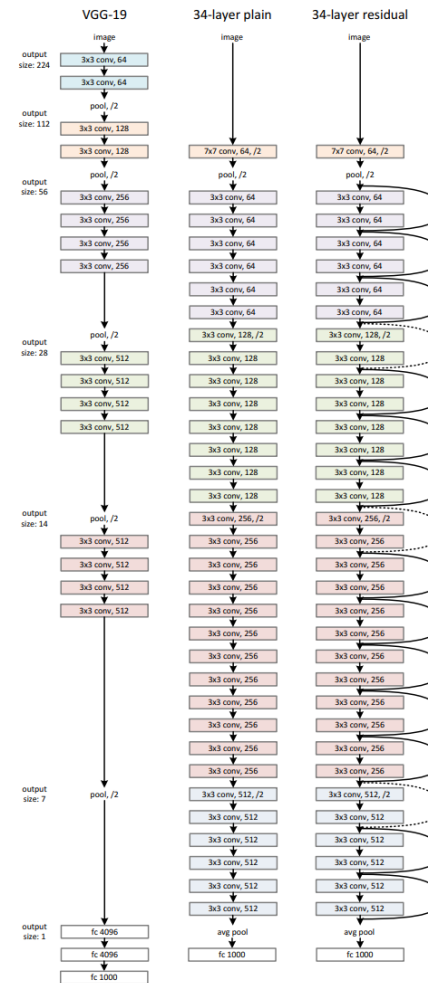
Methods

Model architecture.

CNNs stacked on each other.

Residual connections added in.

VGG nets don't have that.



Methods

They used the following implementation

Images randomly cropped and sampled at 224x224 pixels

Per-pixel mean subtracted

Standard color augmentation

Batch normalization after each convolution

Train from scratch

Use SGD with mini-batch size 256

Learning rate starts at 0.1 and is divided by 10 when the error plateaus

Trained 600,000 iterations with weight decay 0.0001

No dropout

Experiments

Evaluated on ImageNet 2012 classification datasets

- 1000 classes

Evaluated with 18 layer and 34 layer models

Results

18 layer model did better than 34 layers model due to degradation

When residual connections were added, the 34 layer model did better than the 18 layer model, solving degradation.

model	top-1 err.	top-5 err.
VGG-16 [41]	28.07	9.33
GoogLeNet [44]	-	9.15
PReLU-net [13]	24.27	7.38
plain-34	28.54	10.02
ResNet-34 A	25.03	7.76
ResNet-34 B	24.52	7.46
ResNet-34 C	24.19	7.40
ResNet-50	22.85	6.71
ResNet-101	21.75	6.05
ResNet-152	21.43	5.71

Table 3. Error rates (% , **10-crop** testing) on ImageNet validation. VGG-16 is based on our test. ResNet-50/101/152 are of option B that only uses projections for increasing dimensions.

method	top-1 err.	top-5 err.
VGG [41] (ILSVRC'14)	-	8.43 [†]
GoogLeNet [44] (ILSVRC'14)	-	7.89
VGG [41] (v5)	24.4	7.1
PReLU-net [13]	21.59	5.71
BN-inception [16]	21.99	5.81
ResNet-34 B	21.84	5.71
ResNet-34 C	21.53	5.60
ResNet-50	20.74	5.25
ResNet-101	19.87	4.60
ResNet-152	19.38	4.49

Table 4. Error rates (%) of **single-model** results on the ImageNet validation set (except [†] reported on the test set).

method	top-5 err. (test)
VGG [41] (ILSVRC'14)	7.32
GoogLeNet [44] (ILSVRC'14)	6.66
VGG [41] (v5)	6.8
PReLU-net [13]	4.94
BN-inception [16]	4.82
ResNet (ILSVRC'15)	3.57

Table 5. Error rates (%) of **ensembles**. The top-5 error is on the test set of ImageNet and reported by the test server.

Results

101-layer and 152-layer ResNets: We construct 101-layer and 152-layer ResNets by using more 3-layer blocks (Table 1). Remarkably, although the depth is significantly increased, the 152-layer ResNet (11.3 billion FLOPs) still has *lower complexity* than VGG-16/19 nets (15.3/19.6 billion FLOPs).

The 50/101/152-layer ResNets are more accurate than the 34-layer ones by considerable margins (Table 3 and 4). We do not observe the degradation problem and thus enjoy significant accuracy gains from considerably increased depth. The benefits of depth are witnessed for all evaluation metrics (Table 3 and 4).

method			error (%)
Maxout [10]			9.38
NIN [25]			8.81
DSN [24]			8.22
	# layers	# params	
FitNet [35]	19	2.5M	8.39
Highway [42, 43]	19	2.3M	7.54 (7.72±0.16)
Highway [42, 43]	32	1.25M	8.80
ResNet	20	0.27M	8.75
ResNet	32	0.46M	7.51
ResNet	44	0.66M	7.17
ResNet	56	0.85M	6.97
ResNet	110	1.7M	6.43 (6.61±0.16)
ResNet	1202	19.4M	7.93

Summary

ResNets have lower training and validation losses.

ResNets don't overfit as much.

Questions

why was that bad

Provided no useful context for the result

Assumed knowledge the audience didn't have

Focused on implementation details and not the core findings

Didn't provide a critical analysis on results or qualify their impact

Takeaway summary is not actionable advice

Deep Residual Learning for Image Recognition

[2016 IEEE Conference on Computer Vision and Pattern Recognition \(CVPR\)](#)

Kaiming He

Xiangyu Zhang

Shaoqing Ren

Jian Sun

Microsoft Research

{kahe, v-xiangz, v-shren, jiansun}@microsoft.com

Agenda

Today I will present

- (Intro) The problem this paper is trying to solve
- (Methods) The residual connection method
- (Results) The experimental setup in the paper
- (Discussion) Author's interpretation of the findings
- Impact on the field
- Questions

The Problem

The deeper a neural network is, the harder it is to train.

Technically, a model with more parameters should fit better than a model with fewer, right?

Wrong!

The Problem

“Degradation” of performance occurs in deep networks.




Figure 1. Training error (left) and test error (right) on CIFAR-10

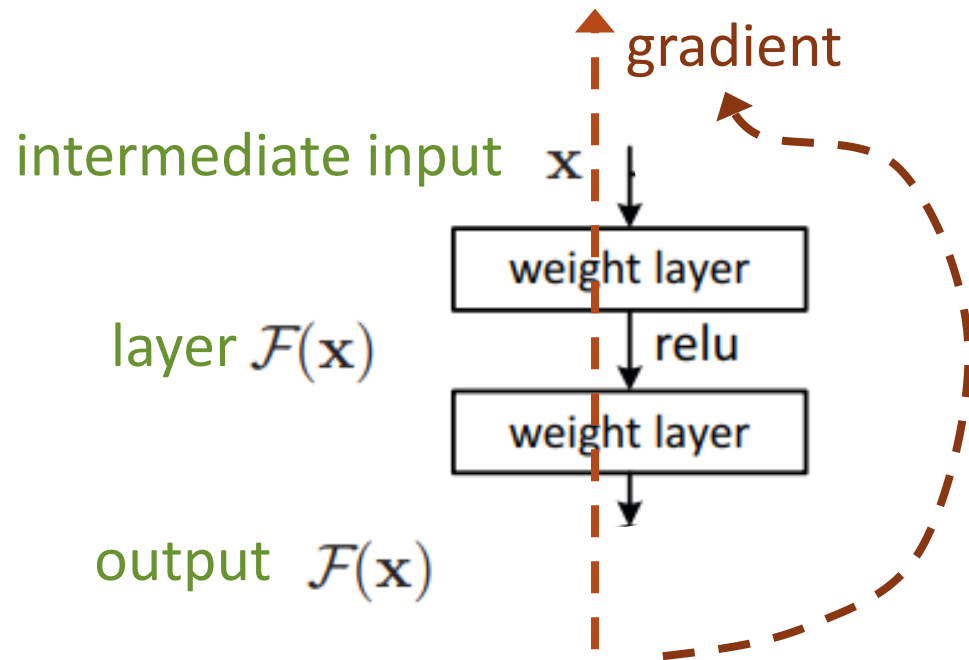
Training a very deep neural network (20+ layers) is always challenging.

Theoretical basis for this challenge is not yet clear. Lots of research has gone into studying this problem, including:

- Batch normalization - <https://arxiv.org/pdf/1502.03167v3.pdf>
- ReLU activation - https://proceedings.neurips.cc/paper_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf
- Weights initialization - <https://proceedings.mlr.press/v9/glorot10a/glorot10a.pdf>

The Method

Authors propose a mechanism to mitigate degradation, called “residual connections”



In this way, a multi-layer block trains like a single layer.

Note: a “residual” connection *adds* the input to the output as it pertains to this paper only. A “skip” connection is a more general form of this that skips multiple layers.

Hypotheses

Deeper models will degrade in performance compared to less deep models in identical setups

Addition of residual connections alone will improve benchmark performance on a deep model

Models trained with residual connections will outperform larger state of the art models

Experiments

Use the ImageNet dataset

- 14 million images
- 100,000 “classes” or “synsets”

Benchmark task is to train on a subset and learn to classify images by the objects present in them.




Experiments

First, identify if degradation does indeed occur with deeper models.

Train two deep convolutional neural net models, an 18-layer and 34-layer.

Does the 34 layer model perform worse?




Experiments


The authors initially set up 2 models for comparison.

1. plain CNN (either 18 or 34 layers deep)
2. residual CNN (18 or 34, identical with addition of residual connections)

34-layer plain



34-layer residual



Experiments


Adding residual connections to deep models improve performance

Top-1 error %

	plain	ResNet
18 layers	27.94	27.88
34 layers	28.54	25.03

Thin = training Top-1 error

Bold = validation Top-1 error



Experiments

How much better can performance get? Is it better than the state of the art model, VGG-16?

<https://arxiv.org/pdf/1409.1556.pdf>

How deep can you go?

model	top-1 err.	top-5 err.
VGG-16 [41]	28.07	9.33
GoogLeNet [44]	-	9.15
PReLU-net [13]	24.27	7.38
plain-34	28.54	10.02
ResNet-34 A	25.03	7.76
ResNet-34 B	24.52	7.46
ResNet-34 C	24.19	7.40
ResNet-50	22.85	6.71
ResNet-101	21.75	6.05
ResNet-152	21.43	5.71

Table 3. Error rates (% , **10-crop** testing) on ImageNet validation. VGG-16 is based on our test. ResNet-50/101/152 are of option B that only uses projections for increasing dimensions.

Interpretation

- Authors assert that degradation impacts deep models, and that residual connections mitigate it. Their evidence is compelling.
- “Degradation” remains a poorly defined concept. Perhaps better characterization of it would help formulate a more principled solution.
- Decisions on *how* to make residual connections was arbitrary.

Impact

Residual connections are an effective trick, but they lack a theoretical underpinning so far.

Top performance on ImageNet benchmark is impressive – but overfitting to a single benchmark dataset can be misleading.

Points us towards better questions to ask like:

- What *exactly* causes degradation?
- Why do residual connections improve deep models ability to learn?
- (engineering) how to choose the kind of residual/skip connection based on the model, data, or problem domain?

Questions

why was that better (hopefully)

Defined the problem that needed solving

Minimal use of jargon, citations where necessary

Did not dwell on the math – presentation time better spent elsewhere

Broke paper down into set of hypotheses and experiments to test them

Provided added value to the paper by lending my own analysis

Thank you!

:)